Software engineering module 1 (Python)

High dimensional data, dimensionality reduction, and manifold learning

with application to single-cell genomics data

Jian Zhou   **TA: Chenlai Shi**

Overview of Module 1

- Mathematical properties of high dimensional space

- Overview of dimensional reduction methods

- Implement t-SNE algorithm

- Implement GraphDR algorithm

- Elements of scientific computing performance:

    - Computational complexity of common operations

    - Ideas for designing faster algorithm (Barnes-Hut t-SNE, FIt-SNE, randomized SVD)

    - Implementation (BLAS, MKL, GPU, Cython, etc)

- Application to single-cell genomics dataset

- Python code and documentation style. Auto-generation of documentation website with Sphinx

- Anatomy of python scientific software (API, CLI, README, Docs, LICENSE).

- Basics of Git, Github, pull request, & code review

- Demo

Grading for module 1 (20% of total grade)

t-SNE + GraphDR

Code 10% (correctness, efficiency, readability)
**Documentation 5%**
Code review 3%
**Exploratory tasks *2-5%** (optional tasks or additional features / explorations)

*You can receive bonus points from exploratory tasks, but total grade won't exceed 20%

Properties of high dimensional space

d=2

d=3

high d

Distance = r

s

r

s

s

Distance = $r\sqrt{d}$

$$\frac{V_{\text{hypersphere}}}{V_{\text{hypercube}}} = \frac{(d/2)!}{\pi^{d/2}} = \frac{1}{d^{\Theta(d)}}$$

- For KNN classifier, proportion of space within a unit radius hypersphere is increasingly small in high dimensions
- The proportion of space with high probability under an isotropic Gaussian distribution is increasing small in high dimensions

Properties of high dimensional space

**Number of almost orthogonal vectors that fit into d-dimensional space**

Figure 1: Number of almost-orthogonal vectors in $\Re^2, \Re^3, \Re^d$

Random vectors are *almost* orthogonal to each other

$$\cos(\theta) = \frac{\langle x, y \rangle}{\|x\|_2 \|y\|_2}$$

```
[-1, -1,  1,  1,  1,  1,  1,  1,  1,  1, …]
[ 1,  1, -1,  1,  1,  1, -1,  1, -1,  1, …]
```

Proof sketch: With high probability, random vectors are almost orthogonal (exponential with d) with Chernoff bound,
With union bound, the number of vectors that are almost orthogonal with high probability grows exponentially with d too.

The "Curse of dimensionality"

The term comes from Bellman, 1961, on combinatorial optimization problems and used to motivate dynamic programming.

It is now widely used to refer to the phenomenon in many field that difficulties arise with high number of dimensions.

High dimensional data has desirable properties too: so-called "blessings of dimensionality"

Real data is typically confined in a region with low effective dimensionality

Johnson–Lindenstrauss lemma

Given a set X of m points in $\mathbb{R}^N$, we would like to find a mapping $f \, \mathbb{R}^N \to \mathbb{R}^n$ where n << N, so that

$$(1 - \varepsilon)\|u - v\|^2 \leq \|f(u) - f(v)\|^2 \leq (1 + \varepsilon)\|u - v\|^2 \quad \text{, for any } u, v \in X$$

In other words, the distance between all pairs of points in preserved.

Johnson–Lindenstrauss lemma: $n = \mathcal{O}\left(\dfrac{\log m}{\epsilon^2}\right)$ dimensions is sufficient, and the mapping can be a linear mapping.

Sketch of proof: consider **random projection (e.g. random sign, random Gaussian projection)** from $\mathbb{R}^N$, such that the **expectation** of $\|f(u) - f(v)\|^2$ equals to $\|u - v\|^2$, then show that the value of $\|f(u) - f(v)\|^2$ **is close to** $\|u - v\|^2$ with sufficiently high probability (note that this probability does not depend on N).

For proof, find https://www.cs.princeton.edu/courses/archive/fall16/cos521/Lectures/lec9.pdf

Remark 1: this applies to any data in $\mathbb{R}^N$. In other words - this is the worst case scenario of any data
Remark 2: this proof only applies to L2 distance, and it has been proven that L1 distance cannot be preserved

Improve over Random Projection with data-dependent linear projection

Problem 1:

Consider low dimensional embedding Z of the high-dimension data X, so that

$$\underset{Z}{\text{minimize}} \sum_{ij} \left(\|X_i - X_j\| - \|Z_i - Z_j\| \right)^2$$
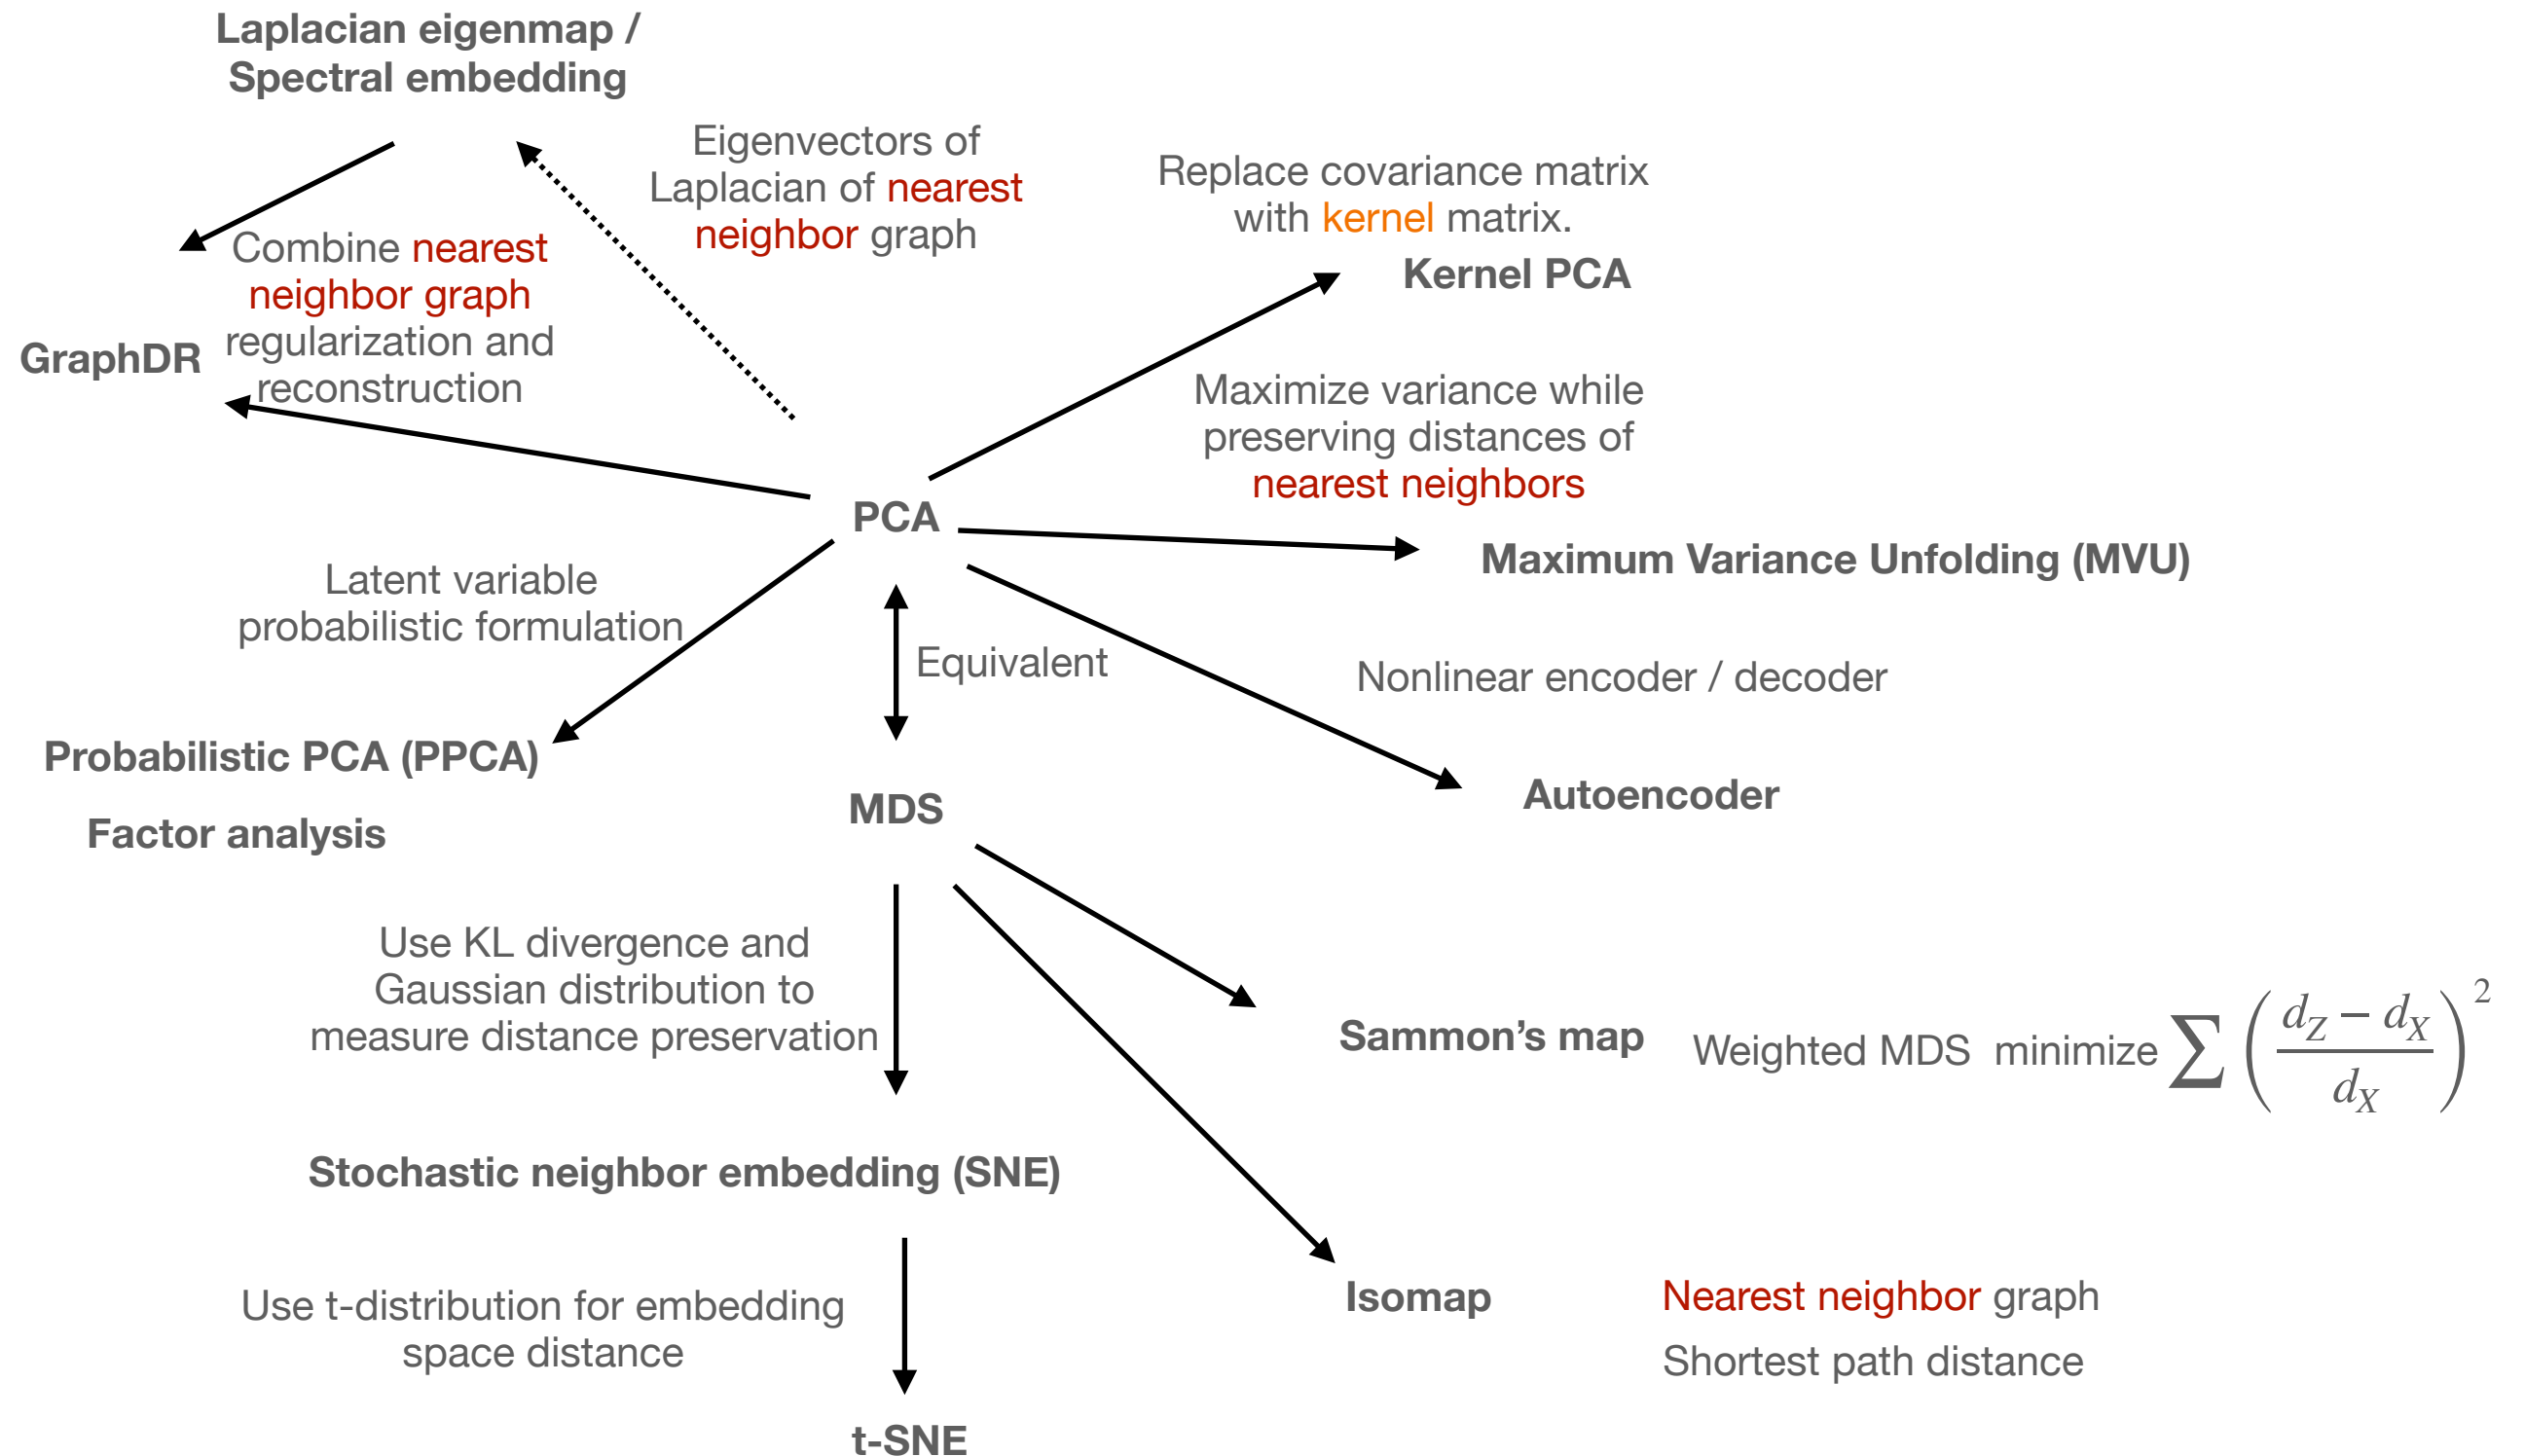
Problem 2:

Consider low dimensional embedding Z of the high-dimension data X, which we express as $Z = XW$,

$$\underset{W}{\text{minimize}} \quad \|X - ZW^T\|^2 \quad \text{subject to} \quad W^T W = I$$

Both formulations (often called classical multidimensional scaling or principal component analysis respectively) are equivalent and lead to the same solution:

A map of dimensionality reduction methods*

**Laplacian eigenmap /
Spectral embedding**

Eigenvectors of
Laplacian of nearest
neighbor graph

Replace covariance matrix
with kernel matrix.

Kernel PCA

Combine nearest
neighbor graph
GraphDR regularization and
reconstruction

Maximize variance while
preserving distances of
nearest neighbors

PCA

Maximum Variance Unfolding (MVU)

Latent variable
probabilistic formulation

Equivalent

Nonlinear encoder / decoder

Probabilistic PCA (PPCA)

Factor analysis

MDS

Autoencoder

Use KL divergence and
Gaussian distribution to
measure distance preservation

Sammon's map   Weighted MDS minimize $\sum \left(\dfrac{d_Z - d_X}{d_X} \right)^2$

Stochastic neighbor embedding (SNE)

Use t-distribution for embedding
space distance

Isomap   Nearest neighbor graph

Shortest path distance

t-SNE

* Note there is more than one way to draw connections between methods and often more than one way to interpret a method
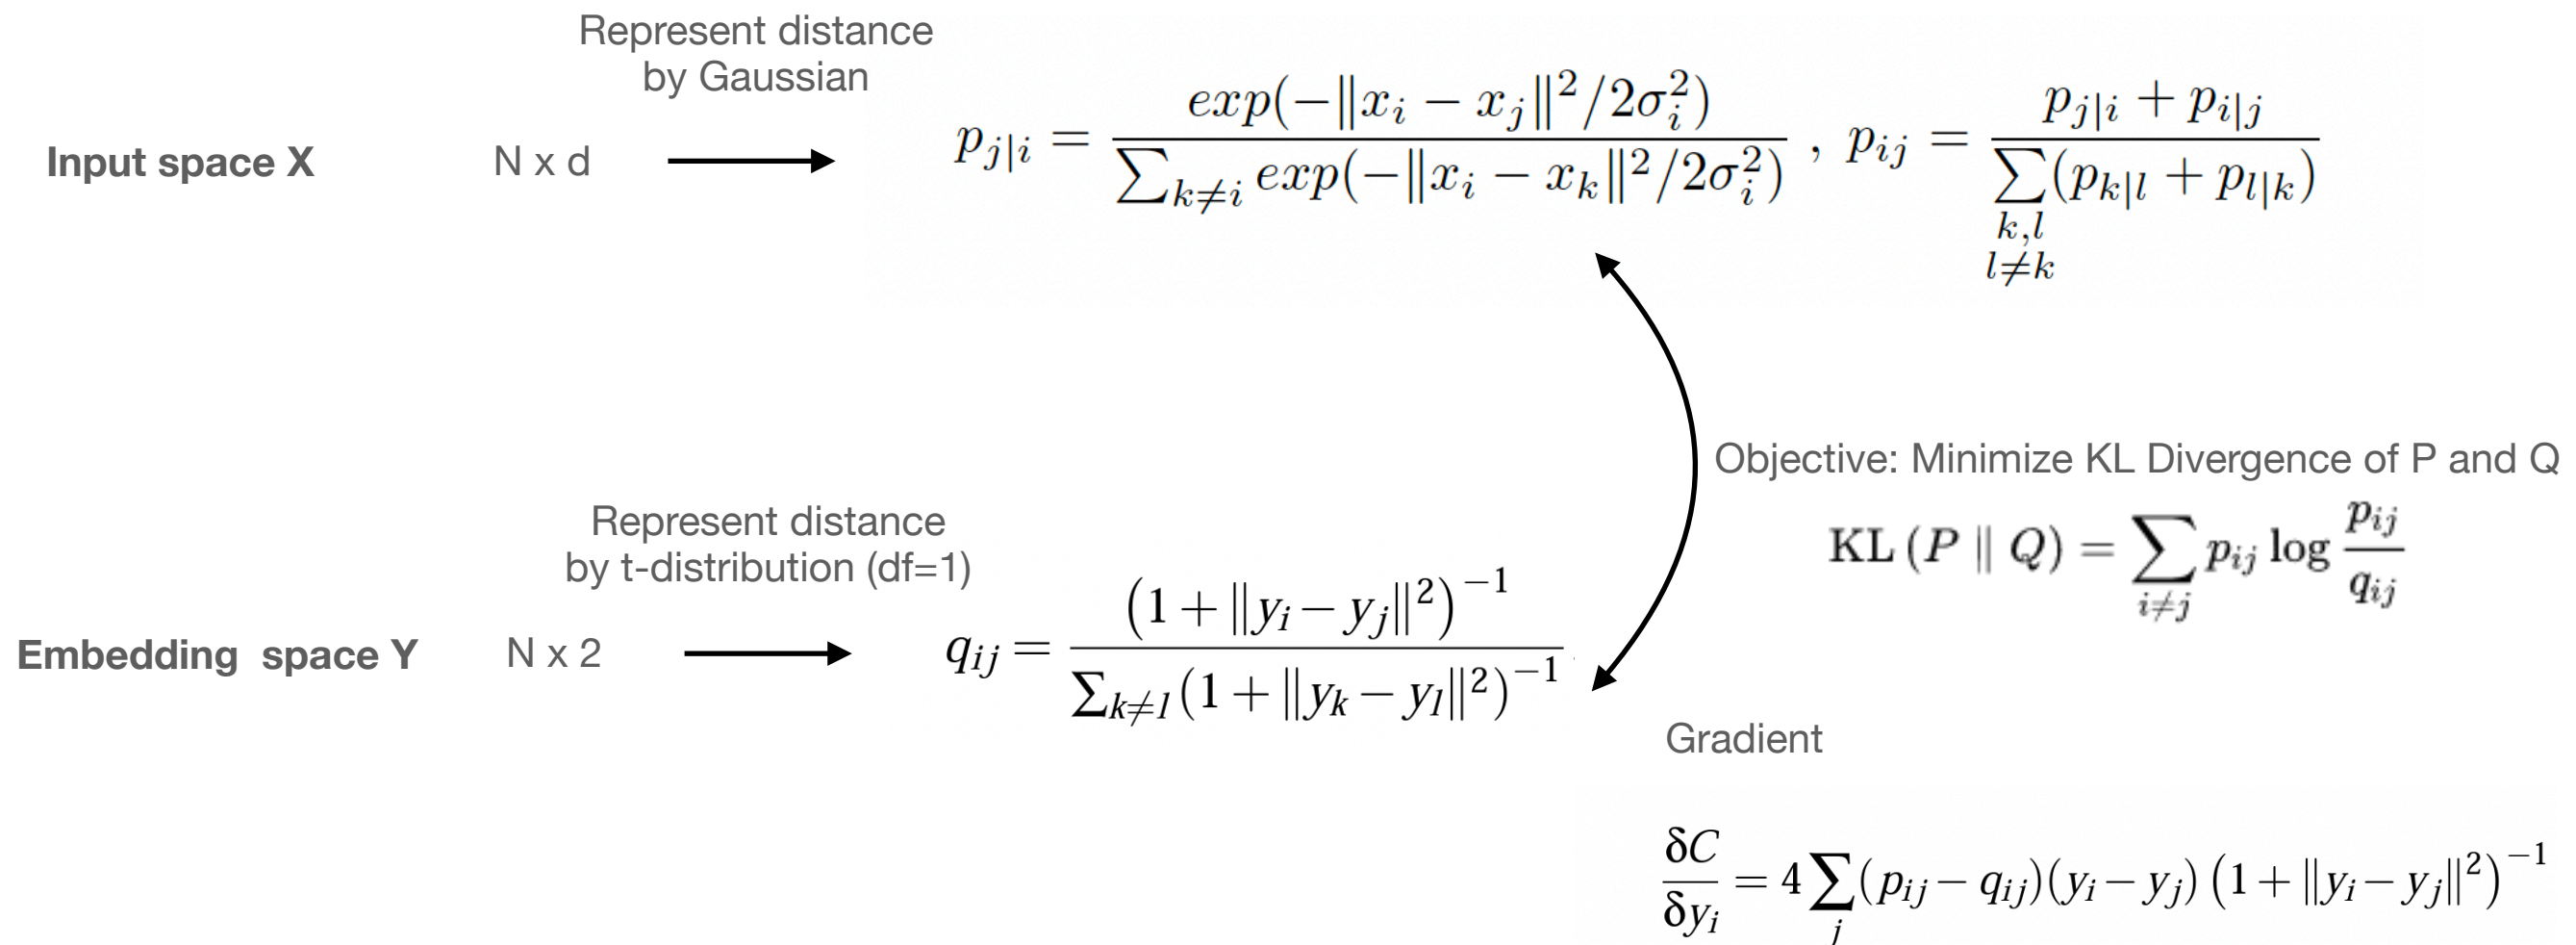
Comparison of t-SNE, Isomap and Sammon's map

t-SNE   **Isomap**   **Sammon's map**

t-SNE (t-distributed stochastic neighbor embedding)

Input space X   N x d

Represent distance
by Gaussian

$$p_{j|i} = \frac{exp(-\|x_i - x_j\|^2/2\sigma_i^2)}{\sum_{k \neq i} exp(-\|x_i - x_k\|^2/2\sigma_i^2)} \;,\; p_{ij} = \frac{p_{j|i} + p_{i|j}}{\sum_{\substack{k,l \\ l \neq k}}(p_{k|l} + p_{l|k})}$$

Objective: Minimize KL Divergence of P and Q

$$\text{KL}\,(P \parallel Q) = \sum_{i \neq j} p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

Embedding space Y   N x 2

Represent distance
by t-distribution (df=1)

$$q_{ij} = \frac{\left(1 + \|y_i - y_j\|^2\right)^{-1}}{\sum_{k \neq l}\left(1 + \|y_k - y_l\|^2\right)^{-1}}$$

Gradient

$$\frac{\delta C}{\delta y_i} = 4 \sum_j (p_{ij} - q_{ij})(y_i - y_j)\left(1 + \|y_i - y_j\|^2\right)^{-1}$$

Comparison of Gaussian
and t-distribution

By IkamusumeFan - Own work, CC BY-SA 3.0, https://commons.wikimedia.org/w/index.php?curid=27359482

t-distribution is heavy-tailed
(Quadratic rather than exponential decay)

t-SNE algorithm breakdown

Loss function:

$$\mathrm{KL}\left(P \parallel Q\right) = \sum_{i \neq j} p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

$$\frac{\delta C}{\delta y_i} = 4 \sum_j (p_{ij} - q_{ij})(y_i - y_j)\left(1 + \|y_i - y_j\|^2\right)^{-1}$$

Optimization:

Gradient descent

$$Y_t = Y_{t-1} + \text{step size} \cdot g_t$$

Gradient descent+momentum

$$Y_t = Y_{t-1} + \Delta Y_t$$

$$\Delta Y_t = \text{momentum} \cdot \Delta Y_{t-1} + \text{step size} \cdot g_t$$

Gradient descent+momentum+
dynamically adjust step size

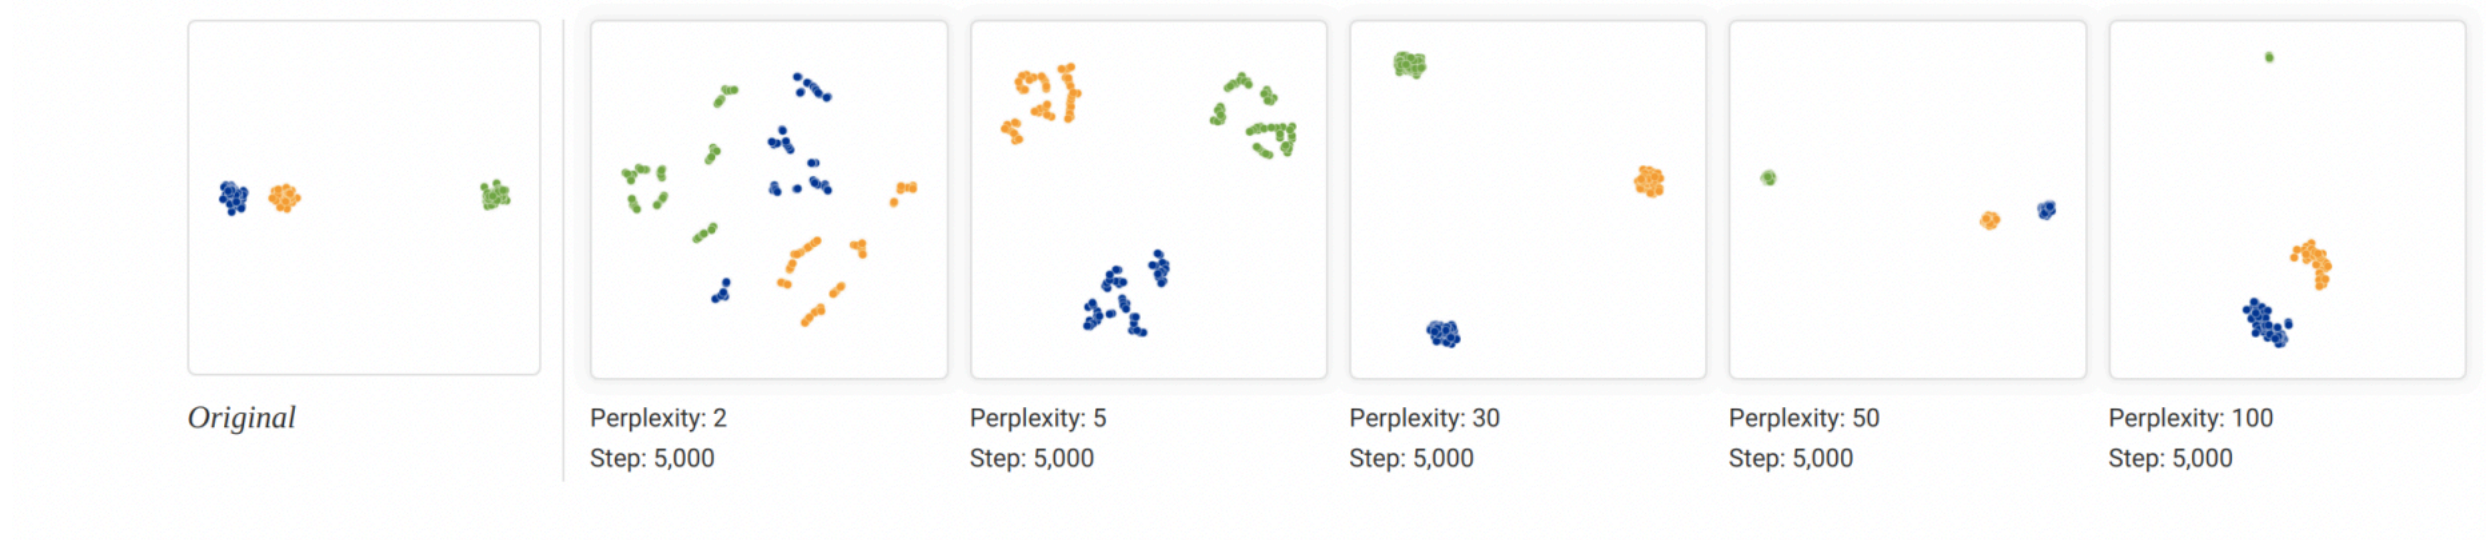|	
Increase step size	If momentum and g signs agrees
Decrease step size	If momentum and g signs disagrees

Early Exaggeration

Multiply p_{ij} by a factor in early iterations,
increase attractive / repulsive force ratio to allow points to
move more freely in early iterations

t-SNE implementation: choosing σ of Gaussian distribution

For each data point, choose a σ so that for the probability distribution $p_{j|i}$ computed from the distances has

entropy $\qquad H = -\sum_{j} p_{j|i} \log p_{j|i} \qquad \approx \log(\text{Perplexity})$

Why is perplexity considered the size of neighbors?

| Original | Perplexity: 2 Step: 5,000 | Perplexity: 5 Step: 5,000 | Perplexity: 30 Step: 5,000 | Perplexity: 50 Step: 5,000 | Perplexity: 100 Step: 5,000 |

We provide the code for choosing $\beta = \dfrac{1}{2\sigma^2}$

t-SNE algorithm pseudocode

Algorithm 1 t-SNE

Input: Data Array $X^{(n \times d)}$
Parameters: no_dims=2, perplexity=30.0, initial_momentum=0.5, final_momentum=0.8, η=500, min_gain=0.01, T=1000
Output: low-dimensional data representation Array $Y^{(n \times no\_dims)}$

Precision(β) adjustment based on perplexity (algorithm 2 and code)
Compute pairwise affinities p_{ij} (equation 1 and note 5.1)
Early exaggerate (multiply) $p_{ij}^{(n \times n)}$ by 4 and clip the value to be at least 10^{-12}
Initialize low-dimensional data representation Array $Y^{(0)}$ using first no_dims of PCs from PCA (code)
Initialize $\Delta Y^{(n \times no\_dims)} = 0$, gains$^{(n \times no\_dims)} = 1$
for $t = 1$ **to** T **do**
    Compute low-dimensional affinities q_{ij} (equation 2 and note 5.2) and clip the value to be at least 10^{-12}
    Compute gradient dY (equation 3 and note 5.3)
    if t < 20 **then**
        momentum = initial_momentum
    else
        momentum = final_momentum
    end if
    Determine gains based on the sign of dY and ΔY :
    gains = (gains + 0.2) * ((dY > 0) \neq (\Delta Y > 0)) + (gains * 0.8) * ((dY > 0.) == (\Delta Y > 0.))
    Clip gains to be at least min_gain
    ΔY = momentum $* \Delta Y - \eta * ($gains $* dY)$
    $Y = Y + \Delta Y$
    if t $== 100$ **then**
        Remove early exaggeration via dividing $p_{ij}^{(n \times n)}$ by 4
    end if
end for

Peformance tip: use numpy broadcasting and array operation whenever you can!

Task for Day 1:

1. **Implement t-SNE following provided pseudocode and guide**
2. **Verify that it is working appropriately on the given MNIST example dataset** (code for loading the dataset and preprocess by PCA provided)
3. Optional: how well does t-SNE work for 3D embedding and can you improve it?

You can work together and help each other, but everyone has to write their own code.
You can ask questions on teams in the Module 1 channel.

Tips for implementation and debugging

Break it down into parts, check your code as you implement part by part

Inspect your intermediate variables (array shape and values)

Use a test dataset (here we use MNIST as in the initial t-SNE paper)

Use an interactive computing environment like Jupyter notebook / Colab or/and an IDE (note that you need to write a .py module file not notebook)

Python coding style - PEP8

Write readable code, avoid bugs, and make collaboration easier (you will review each other's code)

Readings:

PEP8 - Official style guide for python code: https://peps.python.org/pep-0008/
Google python style guide: https://google.github.io/styleguide/pyguide.html

Give your variables meaningful and understandable names

How to start coding first:

Use a formatter (autopep8, yapf, black)

```
pip install black
black mycode.py
```

Use a linter (e.g. pylint)

Python coding style - PEP20

```
>>> import this
The Zen of Python, by Tim Peters

Beautiful is better than ugly.
Explicit is better than implicit.
Simple is better than complex.
Complex is better than complicated.
Flat is better than nested.
Sparse is better than dense.
Readability counts.
Special cases aren't special enough to break the rules.
Although practicality beats purity.
Errors should never pass silently.
Unless explicitly silenced.
In the face of ambiguity, refuse the temptation to guess.
There should be one-- and preferably only one --obvious way to do it.
Although that way may not be obvious at first unless you're Dutch.
Now is better than never.
Although never is often better than *right* now.
If the implementation is hard to explain, it's a bad idea.
If the implementation is easy to explain, it may be a good idea.
Namespaces are one honking great idea -- let's do more of those!
```