

Implementation-level efficiency of python scientific software

Both Algorithm and Implementation level difference can leads to >1000x performance difference. So you have to pay attention to both.

Matrix multiplication of two 4096x4096 matrices

Pure Python

55432s (extrapolated)

Numpy Default BLAS/ATLAS (multi-core)

2.79s

Numpy Intel MKL (multi-core)

265 ms.

CUDA (CUBLAS) with PyTorch

21.7ms

Use fast linear algebra routine whenever possible

What if no existing implementation available?

(Ranked from easy to hard, lower to higher performance improvement potential if used correctly)

Numba JIT

Write Cython code which compiles to C

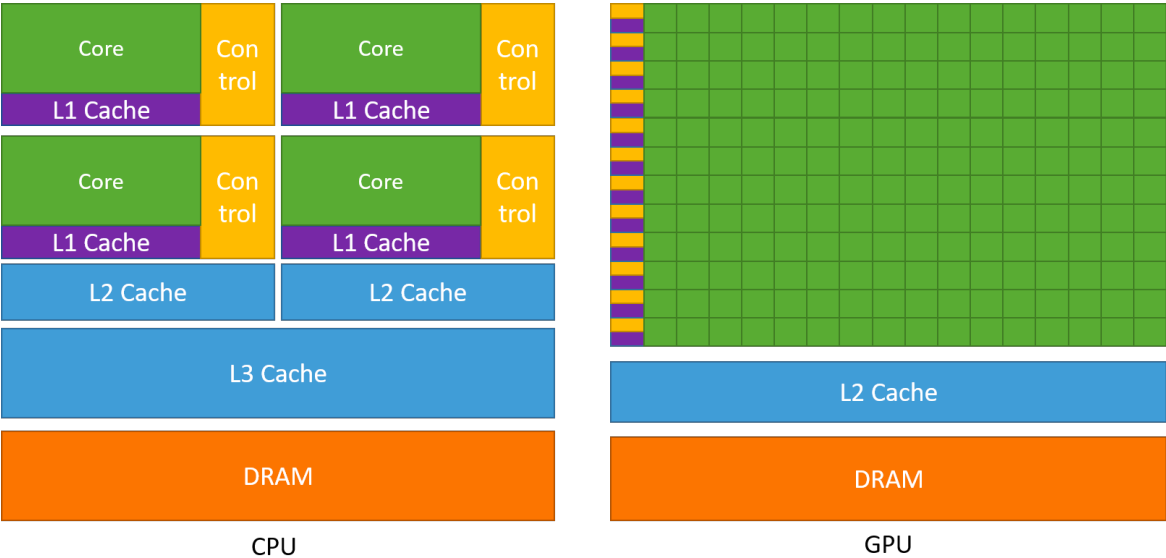
Implement with C/C++ and create pybind11 binding

Low-level optimizations for speed

Built-in parallelism in modern CPU: Single Instruction Multiple Data (SIMD)

Microarchitecture	ISA	throughput (per clock)	SIMD width	max flops/cycle
Nehalem	SSE	1 add + 1 mul	4	8
Sandy Bridge /Ivy Bridge	AVX	1 add + 1 mul	8	16
Haswell	AVX2	2 fmadds	8	32
Knights Corner	AVX-512F	1 fmadd(*)	16	32
Knights Landing	AVX-512F	2 fmadds	16	32

Cache efficiency: communication bandwidth



Cache	Latency	CPU cycles	Size	
	L1 access	~1.2 ns	~4	Between 32 KB and 512 KB
	L2 access	~3 ns	~10	Between 128 KB and 24 MB
	L3 access	~12 ns	~40	Between 2 MB and 32 MB

The order of accessing data matters (improving spatially locality reduces cache miss)

Communication cost for multicore, multi-CPU, or multi-node computation

https://www.eidos.ic.i.u-tokyo.ac.jp/~tau/lecture/parallel_distributed/2016/slides/pdf/simd.pdf

<https://teivah.medium.com/go-and-cpu-caches-af5d32cc5592>

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>

Fast algorithm design!

Fast algorithms:

Barnes-Hut t-SNE

Fast interpolation-based t-SNE

Randomized SVD

Basic ideas:

FFT

Nearest neighbors

Tree

Random projection

Coarse graining


Barnes-Hut t-SNE

$O(N \log N)$.


The objective can be viewed as pairwise attractive and repulsive forces between all cells

$$\frac{\partial C}{\partial \mathbf{y}_i} = 4(F_{attr} + F_{rep}) = 4 \left(\sum_{j \neq i} p_{ij} q_{ij} Z(\mathbf{y}_i - \mathbf{y}_j) - \sum_{j \neq i} q_{ij}^2 Z(\mathbf{y}_i - \mathbf{y}_j) \right)$$

Borrow ideas from N-body simulation!



If points are sufficiently far away,
it is enough to approximate with average of a cell



Fast interpolation-based t-SNE $O(pN + p \log p)$.

The objective can be viewed as pairwise attractive and repulsive forces between all cells

$$\frac{\partial C}{\partial \mathbf{y}_i} = 4(F_{attr} + F_{rep}) = 4 \left(\sum_{j \neq i} p_{ij} q_{ij} Z(\mathbf{y}_i - \mathbf{y}_j) - \sum_{j \neq i} q_{ij}^2 Z(\mathbf{y}_i - \mathbf{y}_j) \right)$$

Borrow ideas from N-body simulation!

Approximate pairwise forces with interpolant with fixed spacing (fast FFT-based acceleration available)

$$\tilde{k}(\mathbf{x}, \mathbf{y}) = \sum_{|\alpha| \leq p} S(\bar{\mathbf{x}}_{\alpha}, \mathbf{x}) \sum_{|\beta| \leq p} \mathbf{k}(\bar{\mathbf{x}}_{\alpha}, \bar{\mathbf{y}}_{\beta}) S(\bar{\mathbf{y}}_{\beta}, \mathbf{y})$$

$O(p^2 + pN)$.

Fast interpolation-based t-SNE $O(pN + p \log p)$.

Further acceleration with FFT (avoiding the p^2 computation):

$k(\bar{x}_\alpha, \bar{y}_\beta)$ between equidistant points can be embedded into a Toeplitz matrix (many relative distances are repeated)

Example Toeplitz matrix

$$\begin{bmatrix} a_0 & a_{-1} & a_{-2} & a_{-3} \\ a_1 & a_0 & a_{-1} & a_{-2} \\ a_2 & a_1 & a_0 & a_{-1} \\ a_3 & a_2 & a_1 & a_0 \end{bmatrix}$$

$$\begin{bmatrix} 7 & 11 & 5 & 6 \\ 3 & 7 & 11 & 5 \\ 8 & 3 & 7 & 11 \\ 1 & 8 & 3 & 7 \end{bmatrix}$$

Toeplitz matrix can be embedded into a circulant matrix like this

$$C_8 = \begin{bmatrix} a_0 & a_{-1} & a_{-2} & a_{-3} & \mathbf{a_0} & \mathbf{a_3} & \mathbf{a_2} & \mathbf{a_1} \\ a_1 & a_0 & a_{-1} & a_{-2} & \mathbf{a_{-3}} & \mathbf{a_0} & \mathbf{a_3} & \mathbf{a_2} \\ a_2 & a_1 & a_0 & a_{-1} & \mathbf{a_{-2}} & \mathbf{a_{-3}} & \mathbf{a_0} & \mathbf{a_3} \\ a_3 & a_2 & a_1 & a_0 & \mathbf{a_{-1}} & \mathbf{a_{-2}} & \mathbf{a_{-3}} & \mathbf{a_0} \\ \mathbf{a_0} & \mathbf{a_3} & \mathbf{a_2} & \mathbf{a_1} & a_0 & a_{-1} & a_{-2} & a_{-3} \\ \mathbf{a_{-3}} & \mathbf{a_0} & \mathbf{a_3} & \mathbf{a_2} & a_1 & a_0 & a_{-1} & a_{-2} \\ \mathbf{a_{-2}} & \mathbf{a_{-3}} & \mathbf{a_0} & \mathbf{a_3} & a_2 & a_1 & a_0 & a_{-1} \\ \mathbf{a_{-1}} & \mathbf{a_{-2}} & \mathbf{a_{-3}} & \mathbf{a_0} & a_3 & a_2 & a_1 & a_0 \end{bmatrix}$$

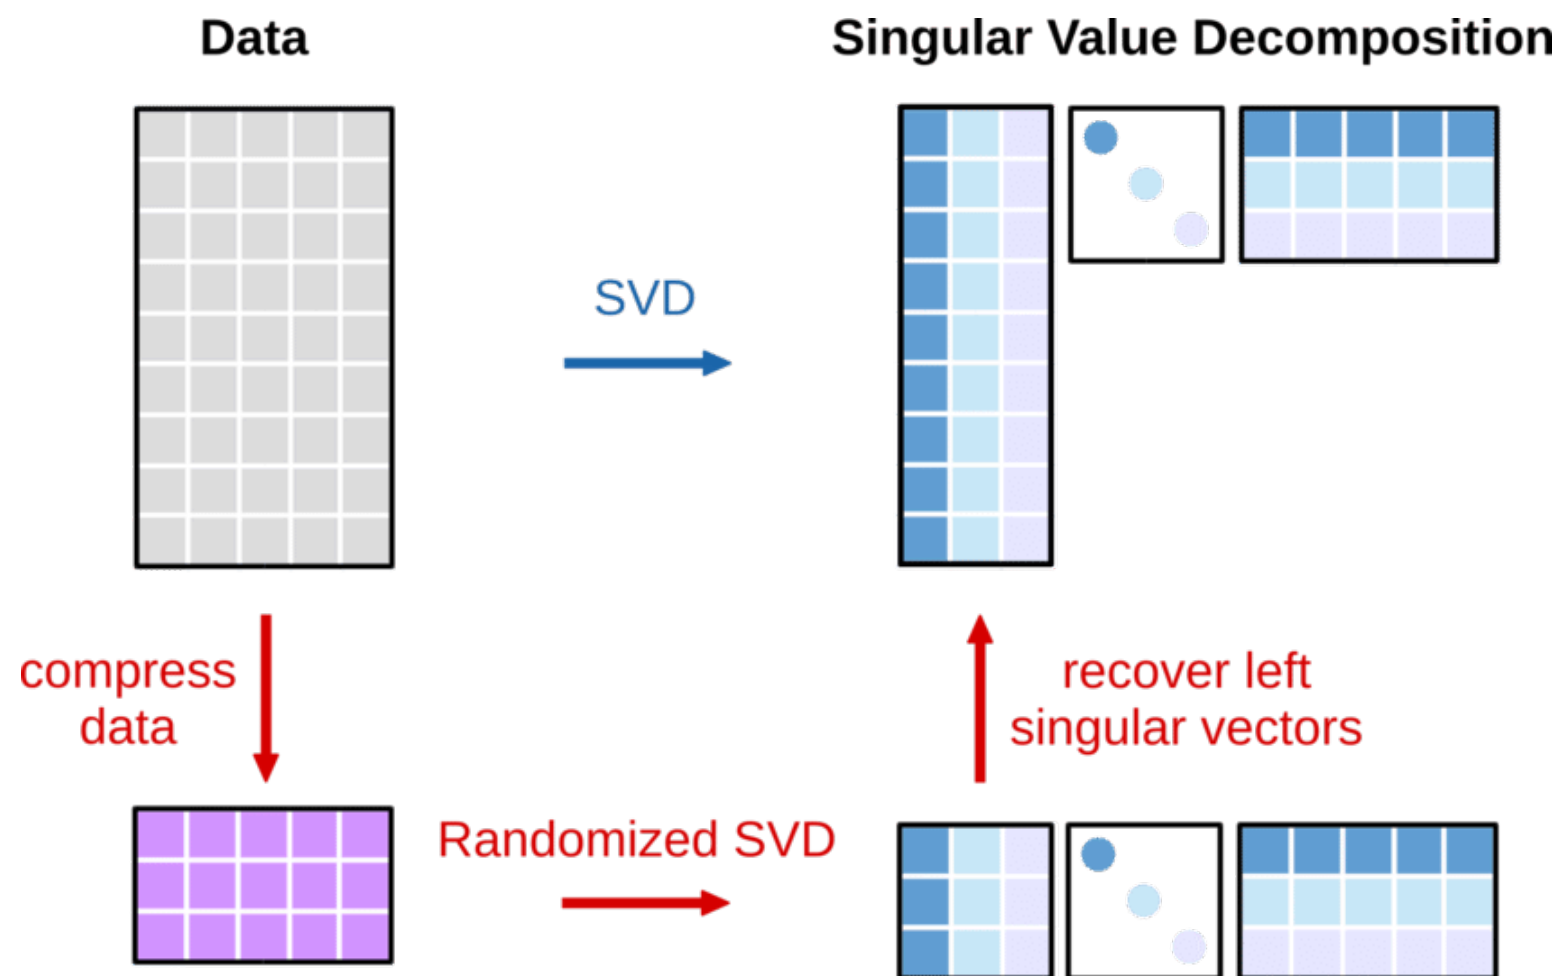
Circulant matrix - vector multiplication can be accelerated with FFT in $O(n \log n)$ time instead of $O(n^2)$

$$C_n x = \text{DFT}^{-1}(\text{DFT}(a_n) \odot \text{DFT}(x))$$

Each DFT or inverse DFT step take $n \log n$ time

Fast algorithm design: Randomized algorithms

Singular value decomposition m-by-n matrix A, $m > n$ $O(mn^2)$




Fast algorithm design: Randomized algorithms

Randomized SVD (sketch)

rank k decomposition for m -by- n matrix A , $m > n$

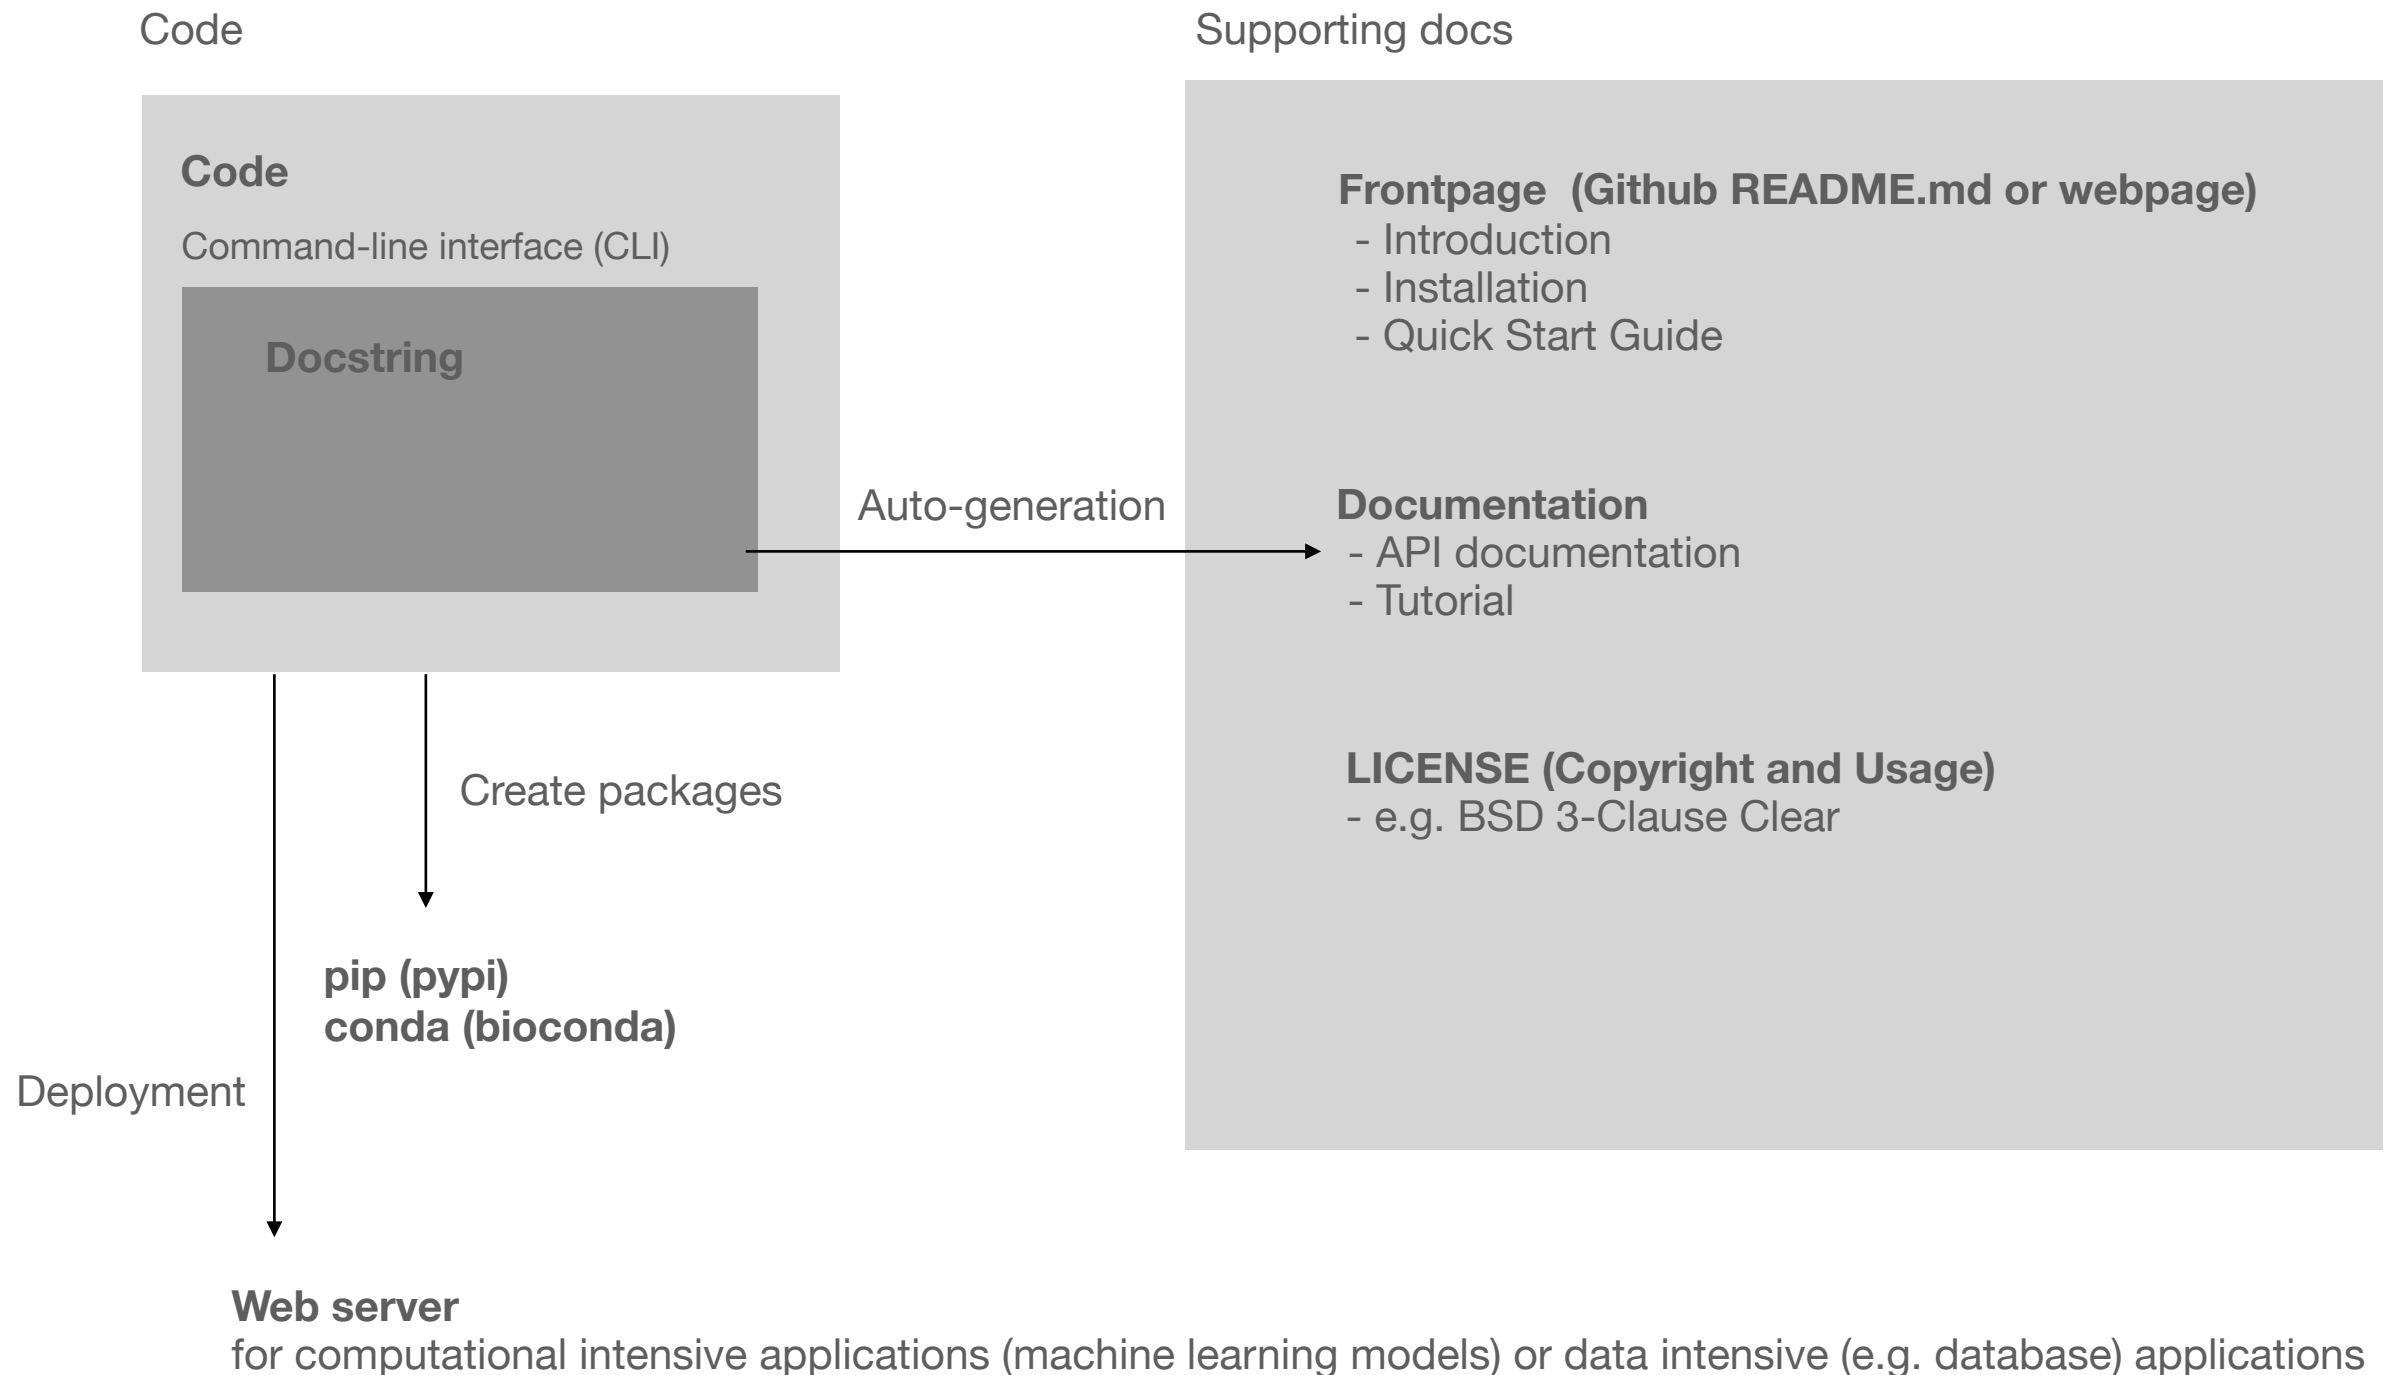
$$O(mnk + mk^2 + nk^2)$$



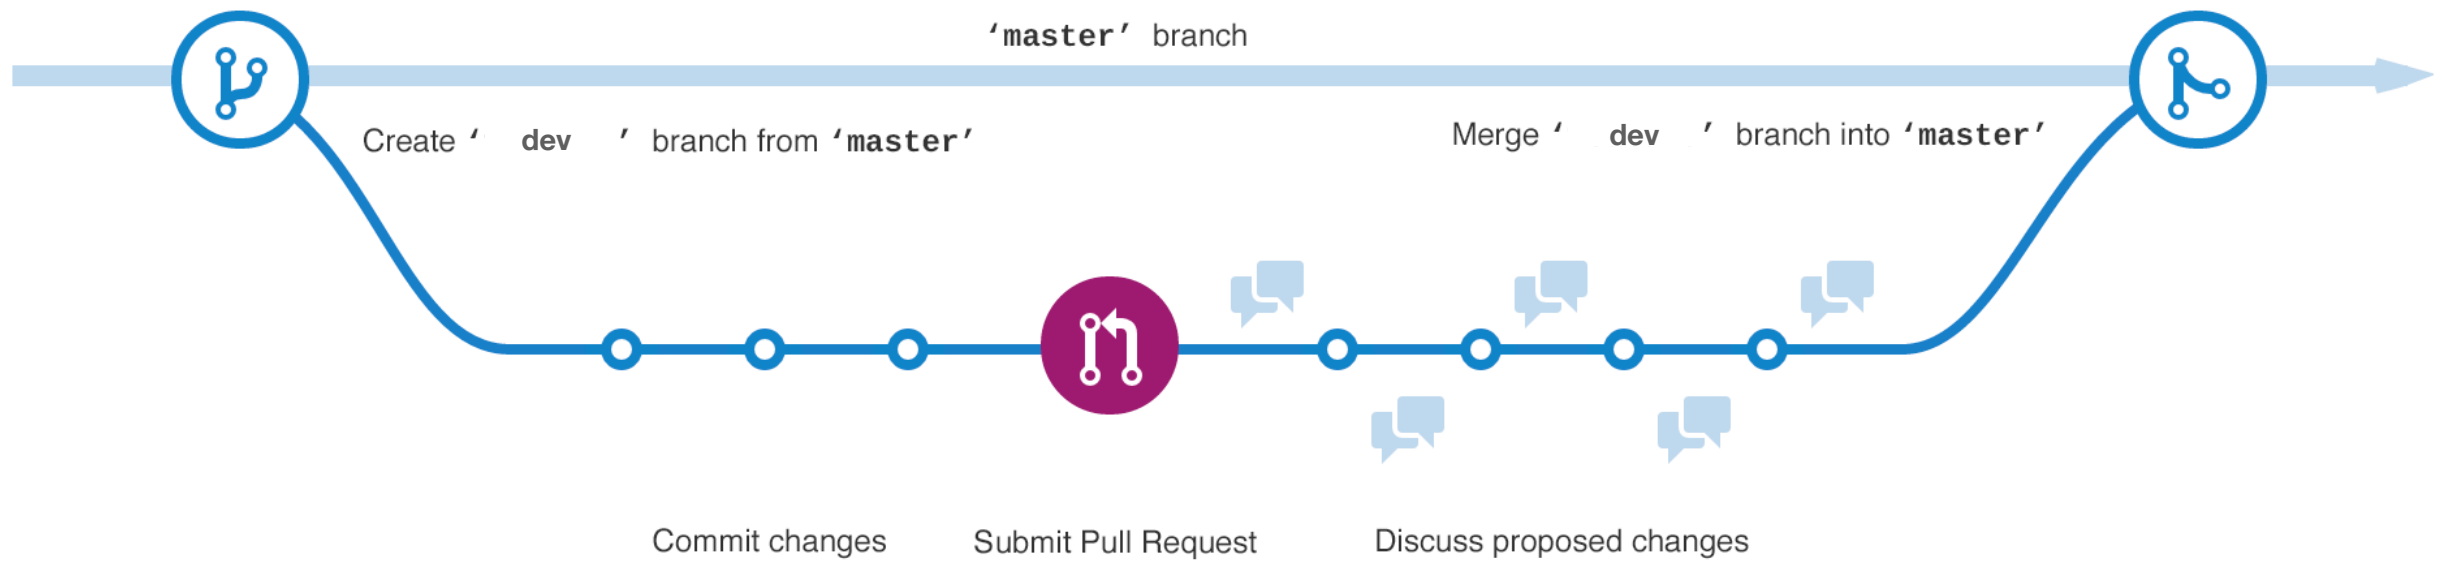
Why?

$$A \approx QQ^T A = Q(\tilde{U}\Sigma V^T) \approx U\Sigma V^T$$

Anatomy of python scientific software



Github workflow



Task for Day 2:

1. Code review (set up Github, create dev branch and pull request)
2. Respond to code review (after approval by your reviewer, merge your code)
3. Write docstring. Optional: implement an CLI with capability to take input file and generate output file / plots; implement scikit-learn Estimator CLI

Docopt : simplest way to make a CLI that parses command line input

```
"""Naval Fate.

Usage:
  naval_fate.py ship new <name>...
  naval_fate.py ship <name> move <x> <y> [--speed=<kn>]
  naval_fate.py ship shoot <x> <y>
  naval_fate.py mine (set|remove) <x> <y> [--moored | --drifting]
  naval_fate.py (-h | --help)
  naval_fate.py --version

Options:
  -h --help      Show this screen.
  --version      Show version.
  --speed=<kn>   Speed in knots [default: 10].
  --moored       Moored (anchored) mine.
  --drifting     Drifting mine.

"""

from docopt import docopt

if __name__ == '__main__':
    arguments = docopt(__doc__, version='Naval Fate 2.0')
    print(arguments)
```

<https://github.com/jazzband/docopt-ng>

<http://try.docopt.org/>